

AirBnB Sydney Price Forecasting

```
In [129...  
# Packages  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import warnings  
warnings.filterwarnings('ignore')  
# this is to clear the warnings from this page, usually we should leave th
```

```
In [130...  
# Plot settings  
sns.set_context('notebook') # optimises figures for notebook display  
sns.set_style('ticks') # set default plot style  
colours = ['#4E79A7', '#F28E2C', '#E15759', '#76B7B2', '#59A14F',  
           '#EDC949', '#AF7AA1', '#FF9DA7', '#9C755F', '#BAB0AB']  
sns.set_palette(colours) # set custom color scheme  
%matplotlib inline  
plt.rcParams['figure.figsize'] = (9, 6)
```

```
In [131...  
train = pd.read_csv('train.csv')  
test = pd.read_csv('test.csv')
```

EDA

```
In [132...  
#Below function drops the $ and other symbols from price, etc. and renders
```

```
In [133...  
def fixsymbols(d):  
    if d is train:  
        d[['price', 'host_response_rate', 'host_acceptance_rate']] = d[['price',  
                                'price']] = d[['price']].astype(float)  
    elif d is test:  
        d[['host_response_rate', 'host_acceptance_rate']] = d[['host_response_  
    else:  
        print('Invalid data set.')
```

```
In [134...  
fixsymbols(train)  
fixsymbols(test)
```

```
In [135...  
print(train['price'].skew())  
print(train['price'].kurtosis())
```

5.7035509554419335
52.62609840971898

In [136...

```
#Taking the log and YJ of price and comparing these.
```


In [137...

```
from sklearn.preprocessing import PowerTransformer

train['logprice'] = np.log(train['price'])


yjp = PowerTransformer(method='yeo-johnson') # YJ is the default, this func
train['priceYJ'] = yjp.fit_transform(train[['price']])

from statlearning import plot_dists
plot_dists(train[['price', 'logprice', 'priceYJ']])
plt.show()
```



In [138...

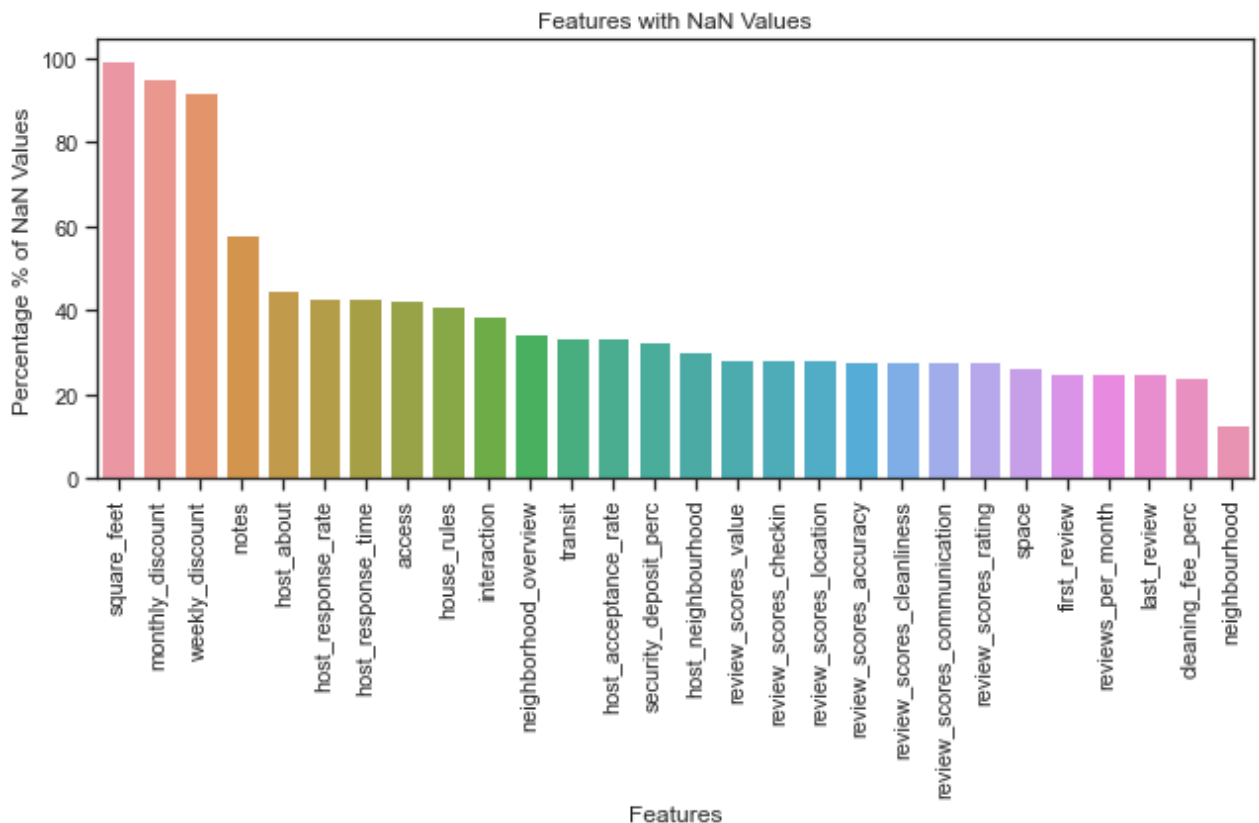
```
g = sns.pairplot(train, kind='reg', x_vars=['latitude', 'longitude'], y_vars='price',
                  plot_kws={'scatter_kws': {'color': sns.color_palette('B', 10)}},
                  g.fig.set_size_inches(13,5))
```



In [139...

```
#Find all features with NaN values
NaN_values = train.isnull().sum().sort_values(ascending=False)[:28]

fig, ax = plt.subplots(1,1, figsize = (9, 6))
sns.barplot(x=NaN_values.index, y=NaN_values/train.shape[0]*100)
plt.xticks(rotation='90')
plt.xlabel('Features')
plt.ylabel('Percentage % of NaN Values')
plt.title('Features with NaN Values')
plt.tight_layout()
plt.savefig('NaNValues.png')
```



In [140...

```
#The below code creates an array imputequantvars which shows all the quant.
#The array imputeobjvars shows all the non-quantitative variables that need
```

In [141...

```
imputevars = train.isnull().sum().sort_values(ascending=False)[train.isnull().sum().sort_values(ascending=False).index[:28]]

imputequantvars = []
imputeobjvars = []
for i in imputevars:
    if train.dtypes[i] != object:
        imputequantvars.append(i)
    if train.dtypes[i] == object:
        imputeobjvars.append(i)
```

Feature Engineering

Data Mining

In [142...

```
# Just an idea
# Best hosts could be considered as the ones that maximise revenue (no.vis.
# There is no visitor count but there is a number of reviews.
# I read somewhere that airbnb thinks that 70% of visits leave a review
# revenue per month of listing as new listings won't have as many reviews
train['revenue_per_month'] = train['price']*(train['reviews_per_month']/0.
train = train[train.revenue_per_month.notna()]
```

Imputing 'object' type variables

In [143...

```
imputeobjvars
#These all need imputing. Below we will make a variable that records wheth
#fills blanks with 'blank'
```

Out[143...

```
['notes',
 'host_about',
 'host_response_rate',
 'host_response_time',
 'access',
 'house_rules',
 'interaction',
 'neighborhood_overview',
 'transit',
 'host_acceptance_rate',
 'host_neighbourhood',
 'space',
 'first_review',
 'last_review',
 'neighbourhood',
 'zipcode',
 'host_location',
 'city']
```

In [144...

```
#HOST RESPONSE RATE is not an object! It is actually an int.

#we will fill it with zeros because a NA response rate probably means no r
train.host_response_rate.fillna(0,inplace=True)
train.host_response_rate = train.host_response_rate.astype(int)

test.host_response_rate.fillna(0,inplace=True)
test.host_response_rate = test.host_response_rate.astype(int)
```

In [145...

```
pd.unique(train.host_response_time)
```

Out[145...

```
array(['within a few hours', 'within an hour', nan, 'within a day',
      'a few days or more'], dtype=object)
```

In [146...

```
#HOST RESPONSE TIME is an ordinal variable

train.host_response_time.fillna(0,inplace=True)
train.host_response_time = np.where(train.host_response_time=='a few days or more',1,train.host_response_time)
train.host_response_time = np.where(train.host_response_time=='within a day',2,train.host_response_time)
train.host_response_time = np.where(train.host_response_time=='within a few days',3,train.host_response_time)
train.host_response_time = np.where(train.host_response_time=='within an hour',4,train.host_response_time)

train.host_response_time = train.host_response_time.astype(int)

#HOST RESPONSE TIME is an ordinal variable

test.host_response_time.fillna(0,inplace=True)
test.host_response_time = np.where(test.host_response_time=='a few days or more',1,test.host_response_time)
test.host_response_time = np.where(test.host_response_time=='within a day',2,test.host_response_time)
test.host_response_time = np.where(test.host_response_time=='within a few days',3,test.host_response_time)
test.host_response_time = np.where(test.host_response_time=='within an hour',4,test.host_response_time)

test.host_response_time = test.host_response_time.astype(int)
```

In [147...

```
train.host_acceptance_rate_blank = np.where(train.host_acceptance_rate.isna(),1,train.host_acceptance_rate)
train.host_acceptance_rate = train.host_acceptance_rate.fillna(0)
train.host_acceptance_rate = train.host_acceptance_rate.astype(int)

test.host_acceptance_rate_blank = np.where(test.host_acceptance_rate.isna(),1,test.host_acceptance_rate)
test.host_acceptance_rate = test.host_acceptance_rate.fillna(0)
test.host_acceptance_rate = test.host_acceptance_rate.astype(int)
```

In [148...

```
#First Review and Last Review are dates. I will make a variable showing the time gap between them
```

In [149...

```
train.first_review.fillna('2000-01-01',inplace=True)
test.first_review.fillna('2000-01-01',inplace=True)
```

In [150...

```
train.first_review = pd.to_datetime(train.first_review)
test.first_review = pd.to_datetime(test.first_review)
```

In [151...

```
train.last_review.fillna('2000-01-01',inplace=True)
test.last_review.fillna('2000-01-01',inplace=True)
```

In [152...

```
train.last_review = pd.to_datetime(train.last_review)
test.last_review = pd.to_datetime(test.last_review)
```

In [153...

```
train.reviewtimegap = train.last_review - train.first_review
test.reviewtimegap = test.last_review - test.first_review
```

```
In [154... train.reviewtimegap=train.reviewtimegap.astype(int)
test.reviewtimegap=test.reviewtimegap.astype(int)
```

```
In [155... #The below function makes a variable for each object variable that still r
#and has NA values.
```

```
In [156... imputeobjvars
```

```
Out[156... ['notes',
'host_about',
'host_response_rate',
'host_response_time',
'access',
'house_rules',
'interaction',
'neighborhood_overview',
'transit',
'host_acceptance_rate',
'host_neighbourhood',
'space',
'first_review',
'last_review',
'neighbourhood',
'zipcode',
'host_location',
'city']
```

```
In [157... imputeobjvars = np.delete(imputeobjvars,[2,3,9,12,13])
```

```
In [158... def imputeobjs(d):
    for i in imputeobjvars:
        if d is train or d is test:
            d[i+'_blank'] = ''
            d[i+'_blank'] = np.where(d[i].isna()==True,1,0)
            d[i] = np.where(d[i].isna()==True,'Blank',d.name.values)
```

```
In [159... imputeobjs(train)
imputeobjs(test)
```

Imputing quantitative variables logically

```
In [160... imputequantvars
```

```
Out[160... ['square_feet',
            'monthly_discount',
            'weekly_discount',
            'security_deposit_perc',
            'review_scores_value',
            'review_scores_checkin',
            'review_scores_location',
            'review_scores_accuracy',
            'review_scores_cleanliness',
            'review_scores_communication',
            'review_scores_rating',
            'reviews_per_month',
            'cleaning_fee_perc',
            'beds',
            'bedrooms',
            'bathrooms']
```

```
In [161... #I have an idea -- maybe places with square feet listed have higher value?
#Let's impute all NAs as 0 and add a dummy for this.

#For all the above quant variables except for beds, bedrooms and bathrooms
```

```
In [162... imputequantvars = imputequantvars[:-3]
```

```
In [163... #Below is a function that will impute all the ones I want.
```

```
In [164... def imputequants(d):
    for i in imputequantvars:
        if d is train or d is test:
            d[i+'_blank'] = np.where(d[i].isna()==True,1,0)
            d[i] = d[i].fillna(0)
            d[i] = d[i].astype(int)
```

```
In [165... imputequants(train)
imputequants(test)
```

Making new dummy variables

Neighbourhood binning

```
In [166... # Fill neighbourhood NaN
train['neighbourhood_cleansed']=train['neighbourhood_cleansed'].fillna(train['neighbourhood_cleansed'])

# This code bins very low frequency neighbourhoods into an 'Other' field -
series = pd.value_counts(train.neighbourhood_cleansed)
```

```
In [167... mask = (series/series.sum() * 100).lt(0.03)
# # To replace df['column'] use np.where I.e
train['neighbourhood_cleansed'] = np.where(train['neighbourhood_cleansed']
```

```
In [168... # Average price of neighbourhood ranked - starting to make sense
# We should bin based on this but need to consider how this will effect the
bins = train.groupby("neighbourhood_cleansed").mean()['price'].sort_values
```

```
In [169... bins
```

```
Out[169... neighbourhood_cleansed
Pittwater      8.0
Mosman         5.0
Hunters Hill   5.0
Manly          5.0
Warringah      5.0
Waverley       4.0
Sutherland Shire 4.0
Lane Cove      4.0
Woollahra      4.0
Leichhardt     4.0
North Sydney   3.0
The Hills Shire 3.0
Randwick       3.0
City Of Kogarah 3.0
Willoughby     3.0
Sydney         3.0
Penrith        3.0
Canada Bay     3.0
Liverpool      3.0
Hurstville     3.0
Hornsby        2.0
Strathfield    2.0
Auburn         2.0
Rockdale       2.0
Fairfield      2.0
Ashfield       2.0
Marrickville   2.0
Camden         2.0
Ku-Ring-Gai    2.0
Ryde           2.0
Parramatta     2.0
Campbelltown   2.0
Canterbury     2.0
Botany Bay     2.0
Bankstown      2.0
Burwood        2.0
Holroyd        1.0
Blacktown      1.0
Name: price, dtype: float64
```


In [170...

```

import matplotlib.pyplot as plt
import cartopy.crs as ccrs
import cartopy.feature as cfeature
from matplotlib.offsetbox import AnchoredText

def main():
    ax = plt.axes(projection=ccrs.PlateCarree())
    ax.set_extent([151.45, 150.6, -34.1, -33.55])

    # Put a background image on for nice sea rendering.
    ax.stock_img()

    # Create a feature for States/Admin 1 regions at 1:50m from Natural Earth
    states_provinces = cfeature.NaturalEarthFeature(
        category='cultural',
        name='admin_1_states_provinces_lines',
        scale='50m',
        facecolor='none')

    SOURCE = 'Natural Earth'
    LICENSE = 'public domain'

    ax.add_feature(cfeature.LAND)
    ax.add_feature(cfeature.COASTLINE)
    ax.add_feature(states_provinces, edgecolor='gray')

    # Get latitude and longitude for each data point
    x, y = (train['longitude'].to_numpy(), train['latitude'].to_numpy())

    # Scatter plot as heat map
    ax.scatter(x, y, c=(train['price']),
              edgecolors='none', cmap=plt.get_cmap('rainbow'), alpha=1)

    sns.scatterplot(x='longitude', y='latitude', hue='price', hue_norm=(0, 4))
    plt.legend([], [], frameon=False)

    # Add a text annotation for the license information to the
    # the bottom right corner.
    text = AnchoredText(r'$\mathcircled{c}$ {}; license: {}'.format(SOURCE, LICENSE),
                        loc=4, prop={'size': 12}, frameon=True)
    ax.add_artist(text)
    plt.title('Airbnb Price Listings in Sydney', fontsize=17, y=1.01, fontv
    plt.show()

if __name__ == '__main__':
    main()

```

Airbnb Price Listings in Sydney



In [171...

```
pd.unique(train.neighbourhood_cleansed)

def neighbourhoodbins(d):
    if d is train or d is test:

        d['neighbourhood_cleansed'] = d['neighbourhood_cleansed'].fillna(d['neighbourhood_cleansed'])
        series = pd.value_counts(d.neighbourhood_cleansed)

        mask = (series / series.sum() * 100).lt(0.03)
        d['neighbourhood_cleansed'] = np.where(d['neighbourhood_cleansed'].isin(mask.index), 0, d['neighbourhood_cleansed'])

        for i in range(len(bins.index)):
            d.neighbourhood_cleansed = np.where(d.neighbourhood_cleansed == bins.index[i], bins.index[i], d.neighbourhood_cleansed)
```

In [172...

```
neighbourhoodbins(train)
neighbourhoodbins(test)

train.neighbourhood_cleansed = train.neighbourhood_cleansed.astype(float)
test.neighbourhood_cleansed = test.neighbourhood_cleansed.astype(float)
```

Here we carefully consider which categorical variables to include as dummies...

What do the most expensive places have in their summary?

In [173...

```
train = pd.get_dummies(train, columns=['cancellation_policy', 'requires_license'])
```

```
In [174... test = pd.get_dummies(test,columns=['cancellation_policy','requires_license
```

Amenity dummies

```
In [175... #Below we define a function for rendering train and test amenity dummies.
```

```
In [176... #This simply adds up all the amenities to generate a score.

def amenityscore(d):

    amenitieslist = []

    for i in range(len(d.amenities.values)):
        k = d.amenities.values[i].replace('{','').replace('}','').replace(
        for j in range(len(k)):
            if k[j] not in amenitieslist:
                amenitieslist.append(k[j])

    if d is train or d is test:
        d['amenityscore'] = 0
        for i in amenitieslist:
            d['amenityscore'] += np.where(d.amenities.str.contains(i,case=
        d['amenityscore'] += np.where(d.summary.str.contains('beach',case=
    else:
        print('Invalid entry, type train or test into the function.')
```

```
In [177... amenityscore(train)
amenityscore(test)
```

```
In [178... #AND/OR make a dummy for each amenity

def amenitify(d):

    amenitieslist = []

    for i in range(len(d.amenities.values)):
        k = d.amenities.values[i].replace('{','').replace('}','').replace(
        for j in range(len(k)):
            if k[j] not in amenitieslist:
                amenitieslist.append(k[j])

    if d is train or d is test:
        for i in amenitieslist:
            d[i] = np.where(d.amenities.str.contains(i,case=False),1,0)
        d['beach'] = np.where(d.summary.str.contains('beach',case=False),1
    else:
        print('Invalid entry, type train or test into the function.')
```

```
In [179... amenitify(train)
amenitify(test)

In [180... train['bedroomsXaccommodates'] = train['bedrooms']*train['accommodates']

In [181... test['bedroomsXaccommodates'] = test['bedrooms']*test['accommodates']

In [182... train['bedroomsXbathrooms'] = train['bedrooms']*train['bathrooms']

In [183... test['bedroomsXbathrooms'] = test['bedrooms']*test['bathrooms']

In [184... train['hotlatitude'] = np.where(((train['latitude']<-33.85)&(train['latitude']>33.85)),1,0)

In [185... test['hotlatitude'] = np.where(((test['latitude']<-33.85)&(test['latitude']>33.85)),1,0)

In [186... train['hotlongitude'] = np.where((train['longitude']>151.2),1,0)

In [187... test['hotlongitude'] = np.where((test['longitude']>151.2),1,0)

In [188... train['hotlatXhotlong'] = train['hotlatitude']*train['hotlongitude']
test['hotlatXhotlong'] = test['hotlatitude']*test['hotlongitude']

In [189... abs(train.corr()['logprice']).sort_values(ascending=False)[3:30].drop('revenue')
```

```

Out[189... accommodates      0.679780
            bedrooms        0.664572
            bedroomsXaccommodates 0.642417
            bedroomsXbathrooms 0.629077
            beds            0.619002
            room_type_Private room 0.580205
            bathrooms       0.428830
            guests_included  0.357145
            neighbourhood_cleansed 0.304222
            Family/kid friendly 0.270347
            TV               0.240464
            amenityscore     0.238698
            cancellation_policy_strict_14_with_grace_period 0.235938
            security_deposit_perc_blank 0.232616
            longitude        0.228388
            calculated_host_listings_count_private_rooms 0.217966
            Crib              0.214264
            Children's books and toys 0.212197
            extra_people_perc 0.208919
            Dishwasher        0.205740
            Indoor fireplace  0.205426
            cleaning_fee_perc_blank 0.203889
            High chair         0.202068
            hotlongitude       0.195786
            Heating            0.190258
            property_type_House 0.188856
            Name: logprice, dtype: float64

```

```

In [190... train.isna().sum().sort_values()

```

```

Out[190... name                0
            Baby bath           0
            Baby monitor        0
            Barbecue utensils   0
            Game console        0
            ..
            bedrooms            3
            bathrooms           3
            bedroomsXaccommodates 3
            bedroomsXbathrooms  6
            beds                22
            Length: 344, dtype: int64

```

```

In [191... features = train.columns.values

```

```

In [192... features = np.delete(features,np.where(features=='priceYJ'))
            features = np.delete(features,np.where(features=='price'))

```

```

In [193... features = np.delete(features,np.where(features=='revenue_per_month'))

```

```
In [194... #AND/OR make a dummy for each verification mode

def verify(d):

    verifylist = []

    for i in range(len(d.host_verifications.values)):
        k = d.host_verifications.values[i].replace("'", '').replace('"', '')
        for j in range(len(k)):
            if k[j] not in verifylist:
                verifylist.append(k[j])
    #return print(verifylist)
    if d is train or d is test:
        for i in verifylist:
            d[i] = np.where(d.host_verifications.str.contains(i, case=False)
        else:
            print('Invalid entry, type train or test into the function.')
```

```
In [195... verify(train)
verify(test)
```

```
In [196... train.host_since = pd.to_datetime(train.host_since)
test.host_since = pd.to_datetime(test.host_since)
```

```
In [197... train['hostlife']=train.host_since.max()-train.host_since
test['hostlife']=test.host_since.max()-test.host_since
```

```
In [198... train.host_since.fillna('2000-01-01', inplace=True)
test.host_since.fillna('2000-01-01', inplace=True)
```

```
In [199... train.hostlife=train.hostlife.astype(int)
test.hostlife=test.hostlife.astype(int)
```

```
In [200... from sklearn.model_selection import train_test_split
X_train, X_val, y_train, y_val = train_test_split(train[features], train['logprice'],
```

```
In [201... fullxtrain = train[np.delete(features, np.where(features=='logprice'))]
fullytrain = train['logprice']
```

```
In [204...
from pandas_profiling import ProfileReport
print('Warning: Pandas Profiling can take several minutes.')
optimise = str(input('Would you like to profile? Answer \'C\' to continue.

if optimise.lower() == 'c':
    profile = ProfileReport(train, minimal=True)
    profile.to_widgets()
else:
    print('Profiling cancelled')
```

Warning: Pandas Profiling can take several minutes.

Profiling cancelled

Modelling

```
In [205...
#One model must be a linear model. We will use an Elastic Net here.
```

```
In [206...
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer

imp = IterativeImputer()
imp.fit(X_train[imputequantvars])
X_train[imputequantvars] = imp.transform(X_train[imputequantvars])
X_val[imputequantvars] = imp.transform(X_val[imputequantvars])
```

Should we be fitting the imputer to the test?

```
In [207...
#imp.fit(test[imputequantvars])
test[imputequantvars] = imp.transform(test[imputequantvars])
```

```
In [208...
fullxtrain[imputequantvars] = imp.fit_transform(fullxtrain[imputequantvars])
```

```
In [209...
#Fixing variable skewness
```

```
In [210...
skews = ['square_feet', 'weekly_discount', 'minimum_nights', 'minimum_minimum_nights']
```

```
In [211...
yj = PowerTransformer(method='yeo-johnson')

X_train[skews] = yj.fit_transform(X_train[skews])
X_val[skews] = yj.transform(X_val[skews])
test[skews] = yj.transform(test[skews])
```

```
In [212... #Dropping training outliers
```

```
In [213... fullxtrain[skews] = yj.fit_transform(fullxtrain[skews])
```

Outliers

```
In [214... quantitativevariables = []
for i in features:
    if train.dtypes[i] == int or train.dtypes[i] == float and i != 'logprice':
        quantitativevariables.append(i)
```

```
In [215... X_train = X_train.fillna(0)
X_val = X_val.fillna(0)
```

```
In [216... fullxtrain = fullxtrain.fillna(0)
```

```
In [217... test = test.fillna(0)
```

```
In [218... from sklearn.linear_model import LinearRegression

ols = LinearRegression()
ols.fit(X_train[quantitativevariables], y_train)

y_fitted = ols.predict(X_train[quantitativevariables])
e = y_train - y_fitted

#e.reset_index(drop=True,inplace=True)

outliers = e.abs().sort_values().iloc[-10:] # retrieves the indexes for the outliers
```

```
In [219... outlierlist = outliers.index.values
```

```
In [220... for i in outlierlist:
    X_train['outlier'] = np.where(X_train.index==i,1,0)
    fullxtrain['outlier'] = np.where(fullxtrain.index==i,1,0)
```

Elastic Net

```
In [221... features = abs(train.corr()['logprice']).sort_values(ascending=False)[3:].index
```



```
In [222... features=np.delete(features,np.where(features=='property_type_Castle')) #cast
features=np.delete(features,np.where(features=='revenue_per_month')) #cast
```

```
In [223... badfeatures = ['Nespresso machine', 'Natural gas barbeque', 'Sonos sound sy
#These features don't exist in the test and so will be deleted.
for i in badfeatures:
    features=np.delete(features,np.where(features==i)) # is causing it to
morebadfeatures = ['zhima_selfie', 'sent_id', 'reviews', 'manual_online
for i in morebadfeatures:
    features=np.delete(features,np.where(features==i)) # is causing it to
```

```
In [224... from sklearn.preprocessing import RobustScaler
scaler = RobustScaler() #scaling

X_train_scaled = scaler.fit_transform(X_train[features])
X_train = pd.DataFrame(X_train_scaled)

X_val_scaled = scaler.transform(X_val[features])
X_val = pd.DataFrame(X_val_scaled)
```

```
In [225... test_scaled = scaler.transform(test[features])
test = pd.DataFrame(test_scaled)
```

```
In [226... fullxtrain = pd.DataFrame(scaler.fit_transform(fullxtrain[features]))
```

```
In [227... from sklearn.linear_model import ElasticNetCV
enet_cv = ElasticNetCV(l1_ratio=[0.01,0.05,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8
enet_cv.fit(X_train,np.ravel(y_train))
```

```
Out[227... ElasticNetCV(cv=5,
              l1_ratio=[0.01, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8,
              0.9,
                      0.95, 0.99, 1])
```

```
In [228... enet_cv.l1_ratio_
```

```
Out[228... 1.0
```

```
In [229... enet_cv.alpha_
```

```
Out[229... 0.0028861641656963383
```

```
In [230... from sklearn.linear_model import ElasticNet
enet = ElasticNet(alpha = enet_cv.alpha_, l1_ratio = enet_cv.l1_ratio_)
```

In [231...

```

from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
kf=KFold(5)
scores = cross_val_score(enet, X_val, y_val, cv=kf, scoring = 'neg_mean_sq

# print the score for each fold
print(np.sqrt(-scores).round(4))
print((sum(np.sqrt(-scores))/len(np.sqrt(-scores))).round(4))
print('RMSE score for Elastic Net model')

[0.3353 0.335  0.3753 0.3212 0.3308]
0.3395
RMSE score for Elastic Net model

```

In [232...

```
enet.fit(X_train,y_train)
```

Out[232...

```
ElasticNet(alpha=0.0028861641656963383, l1_ratio=1.0)
```

In [233...


```

#interpretable fitted and resid
# fitted = np.exp(enet_cv.predict(X_val))
# resid = np.exp(y_val) - np.exp(enet_cv.predict(X_val))
fitted = enet.predict(X_val)
resid = y_val - enet.predict(X_val)

fig, ax= plt.subplots()
sns.regplot(fitted, resid, lowess=True, ax=ax, scatter_kws={'s': 35, 'alpha
ax.set_xlabel('Fitted values', {'fontsize': 12})
ax.set_ylabel('Residuals', {'fontsize': 12})
ax.set_title('Fitted against residuals')
plt.axhline(color='Black', alpha=0.3, linestyle='--')
#EXOGENEITY plot

```

Out[233... <matplotlib.lines.Line2D at 0x7fa75d552a90>



```
In [234... elasticnetvars = pd.DataFrame(enet.coef_.round(3), index = features).T.where
```

```
In [235... enetselect = elasticnetvars.iloc[0][elasticnetvars.iloc[0].values==True].i
```

```
In [236... np.where(elasticnetvars==True)[1]
```

```
Out[236... array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12,
        13, 14, 18, 19, 20, 21, 23, 24, 25, 26, 28, 30, 31,
        32, 35, 40, 41, 42, 44, 45, 46, 47, 48, 54, 55, 62,
        67, 73, 75, 76, 85, 88, 90, 94, 95, 97, 104, 105, 108,
        109, 110, 124, 127, 131, 132, 134, 139, 140, 151, 165, 169, 179,
        190, 198, 205, 237, 239, 240, 254, 277, 278, 291])
```

```
In [237... len(enetselect)
```

Out[237... 75

```
In [238... enetselect
```

```
Out[238... array(['accommodates', 'bedrooms', 'bedroomsXaccommodates',
        'bedroomsXbathrooms', 'beds', 'room_type_Private room',
        'bathrooms', 'guests_included', 'neighbourhood_cleansed',
        'Family/kid friendly', 'TV', 'amenityscore',
        'cancellation_policy_strict_14_with_grace_period',
        'security_deposit_perc_blank', 'longitude', 'extra_people_perc',
        'Dishwasher', 'Indoor fireplace', 'cleaning_fee_perc_blank',
        'hotlongitude', 'Heating', 'property_type_House', 'Cable TV',
        'Air conditioning', 'calculated_host_listings_count_entire_homes',
        'BBQ grill', 'latitude', 'host_listings_count', 'Smoking allowed',
        'Lock on bedroom door', 'Iron', 'hotlatXhotlong', 'Dryer',
        'Coffee maker', 'Pool', 'availability_365',
        'property_type_Apartment', 'Hair dryer', 'Fire extinguisher',
        'hotlatitude', 'house_rules_blank', 'Beachfront',
        'Dishes and silverware', 'First aid kit', 'notes_blank',
        'Microwave', 'host_id', 'Waterfront', 'Refrigerator',
        'host_about_blank', 'Long term stays allowed',
        'instant_bookable_t', 'availability_90',
        'host_neighbourhood_blank', 'Shampoo', 'maximum_minimum_nights',
        'availability_60', 'Free street parking', 'review_scores_rating',
        'review_scores_cleanliness', 'review_scores_location',
        'availability_30', 'weekly_discount', 'review_scores_accuracy',
        'Luggage dropoff allowed', 'number_of_reviews',
        'neighbourhood_blank',
        'calculated_host_listings_count_shared_rooms',
        'number_of_reviews_ltm', 'security_deposit_perc',
        'reviews_per_month', 'Elevator', 'Single level home',
        'cleaning_fee_perc', 'review_scores_value'], dtype=object)
```

```
In [284... enet.fit(fullxtrain,fullytrain)
```

```
Out[284... ElasticNet(alpha=0.0028861641656963383, l1_ratio=1.0)
```

```
In [285... #predictions = yjp.inverse_transform(enet_cv.predict(test).reshape(-1,1))
predictions = np.round(np.exp(enet.predict(test)),2)
```

```
In [286... np.isnan(predictions).sum()
```

```
Out[286... 0
```

```
In [287... np.where(np.isnan(predictions)==True)
```

```
Out[287... (array([], dtype=int64),)
```

```
In [288... my_submission = pd.DataFrame({'id': test.index, 'price': predictions.ravel
my_submission.to_csv('submission.csv', index=False)
```

```
In [243... enet_cv.coef_[enet_cv.coef_!=0]
```

```
Out[243... array([ 1.44101122e-01,  1.33250452e-01, -3.63349653e-02,  7.62526235e-02,
 7.89937619e-03, -4.62849608e-01, -8.67279493e-03,  3.59084154e-03,
 2.79606144e-02,  1.13285401e-02,  3.49333742e-02,  1.38981100e-02,
 5.74739564e-02, -6.25050203e-02,  5.69165200e-02, -3.61235687e-02,
 4.49792524e-02,  3.95572228e-02, -1.51207761e-01,  2.71031709e-02,
 3.70857132e-03,  3.05443821e-02,  3.82751428e-02,  5.61184064e-02,
 2.57758917e-03,  2.06895770e-02,  5.29086845e-03, -9.75963371e-04,
-6.02527570e-17, -2.01244156e-04, -2.69296742e-02, -1.81601914e-02,
 6.27051737e-03,  5.46134718e-02,  2.32771247e-02,  2.15145981e-02,
 3.26408718e-02,  4.75552807e-02, -3.30615022e-03,  8.53249431e-03,
 7.89389960e-04,  3.70736376e-02, -8.78994205e-04,  2.28625361e-02,
-2.09412839e-03,  2.25209107e-02, -2.33080230e-02, -1.80473587e-02,
-4.01637091e-03,  2.11709825e-02, -8.86795927e-03, -7.84073334e-04,
-6.79814700e-03, -1.19586848e-02,  2.44962485e-02, -2.27897114e-02,
 1.58971823e-02,  2.26631857e-02,  9.63520488e-03, -5.21827781e-02,
 1.55270433e-02,  2.12183522e-02,  1.29973704e-02,  1.53444661e-02,
-2.83860097e-04,  2.58545088e-02, -5.06790468e-03, -1.52544624e-02,
-6.02676655e-03, -7.23949640e-02, -1.47508899e-02, -1.16252369e-02,
-1.40113283e-02, -1.62804439e-02,  2.58939789e-02, -1.55885321e-04,
-6.43101485e-03, -2.10552470e-01, -4.16602526e-02])
```

Ordinary Least Squares

We will take the same variables selected above as the L1 ratio = 1, meaning that LASSO has selected these variables and not Ridge.

```
In [244... OLSfeatures = np.where(elasticnetvars==True)[1]
```

```
In [245... ols = LinearRegression()
ols.fit(X_train[OLSfeatures], y_train)
```

```
Out[245... LinearRegression()
```

```
In [246... kf=KFold(5)
scores = cross_val_score(ols, X_val[OLSfeatures], y_val, cv=kf, scoring =

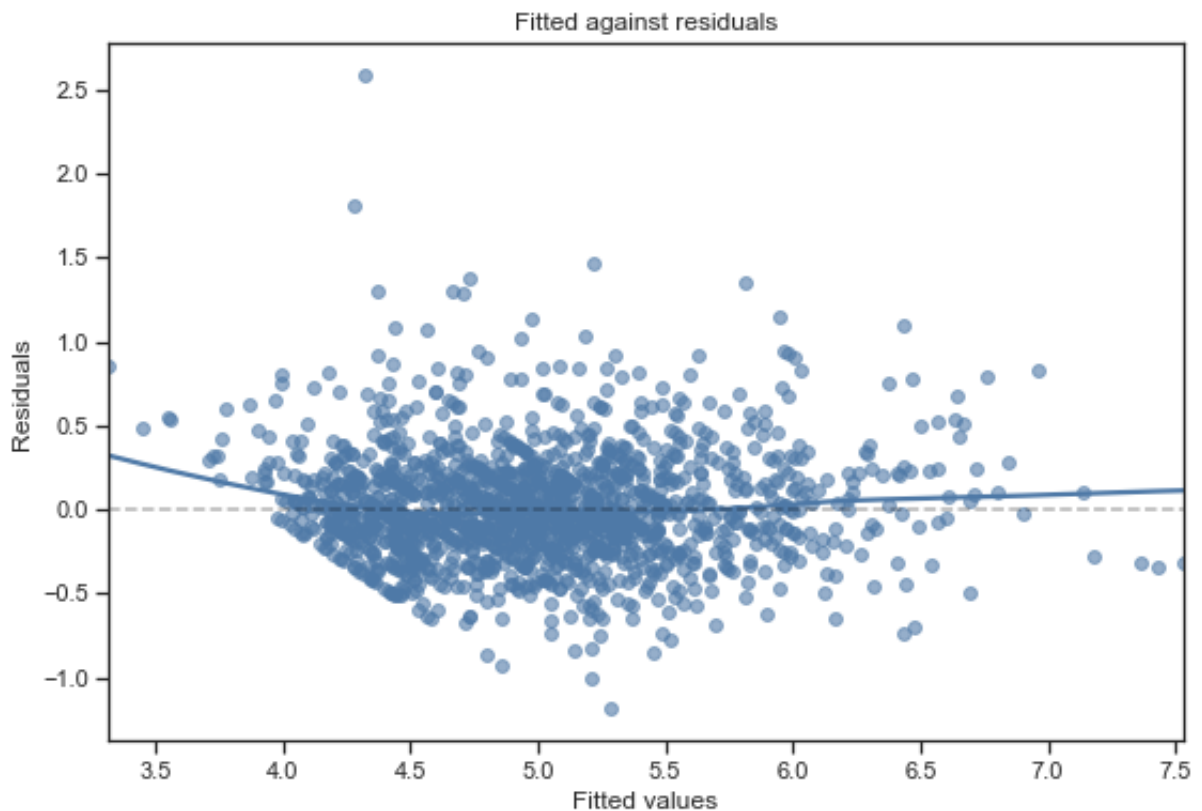
# print the score for each fold
print(np.sqrt(-scores).round(4))
print((sum(np.sqrt(-scores))/len(np.sqrt(-scores))).round(4))
print('RMSE score for Ordinary Least Squares model')

[0.3428 0.3355 0.3764 0.3237 0.3318]
0.342
RMSE score for Ordinary Least Squares model
```

```
In [247... fitted = ols.predict(X_val[OLSfeatures])
resid = y_val - ols.predict(X_val[OLSfeatures])

fig, ax= plt.subplots()
sns.regplot(fitted, resid, lowess=True, ax=ax, scatter_kws={'s': 35, 'alpha': 0.3})
ax.set_xlabel('Fitted values', {'fontsize': 12})
ax.set_ylabel('Residuals', {'fontsize': 12})
ax.set_title('Fitted against residuals')
plt.axhline(color='Black', alpha=0.3, linestyle='--')
#EXOGENEITY plot
```

Out[247... <matplotlib.lines.Line2D at 0x7fa75d552650>



Ridge

We will take the same variables selected before as the L1 ratio = 1, meaning that LASSO has selected these variables and not Ridge.

```
In [248... OLSfeatures = np.where(elasticnetvars==True)[1]
```

```
In [249... from sklearn.linear_model import RidgeCV
ridge_cv = RidgeCV(alphas=np.arange(150,300)*0.1, cv=5)
ridge_cv.fit(X_train[OLSfeatures], np.ravel(y_train))
```

```
Out[249...] RidgeCV(alphas=array([15. , 15.1, 15.2, 15.3, 15.4, 15.5, 15.6, 15.7, 15.8,
15.9, 16. ,
16.1, 16.2, 16.3, 16.4, 16.5, 16.6, 16.7, 16.8, 16.9, 17. , 17.1,
17.2, 17.3, 17.4, 17.5, 17.6, 17.7, 17.8, 17.9, 18. , 18.1, 18.2,
18.3, 18.4, 18.5, 18.6, 18.7, 18.8, 18.9, 19. , 19.1, 19.2, 19.3,
19.4, 19.5, 19.6, 19.7, 19.8, 19.9, 20. , 20.1, 20.2, 20.3, 20.4,
20.5, 20.6, 20.7, 20.8, 20.9, 21. , 21.1, 21.2, 21.3, 21.4, 21.5,
21.6,...
22.7, 22.8, 22.9, 23. , 23.1, 23.2, 23.3, 23.4, 23.5, 23.6, 23.7,
23.8, 23.9, 24. , 24.1, 24.2, 24.3, 24.4, 24.5, 24.6, 24.7, 24.8,
24.9, 25. , 25.1, 25.2, 25.3, 25.4, 25.5, 25.6, 25.7, 25.8, 25.9,
26. , 26.1, 26.2, 26.3, 26.4, 26.5, 26.6, 26.7, 26.8, 26.9, 27. ,
27.1, 27.2, 27.3, 27.4, 27.5, 27.6, 27.7, 27.8, 27.9, 28. , 28.1,
28.2, 28.3, 28.4, 28.5, 28.6, 28.7, 28.8, 28.9, 29. , 29.1, 29.2,
29.3, 29.4, 29.5, 29.6, 29.7, 29.8, 29.9])),
cv=5)
```

```
In [250...] ridge_cv.alpha_
```

```
Out[250...] 22.700000000000003
```

```
In [251...] from sklearn.linear_model import Ridge
ridge = Ridge(alpha=ridge_cv.alpha_)
```

```
In [252...] from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
kf=KFold(5)
scores = cross_val_score(ridge, X_val[OLSfeatures], y_val, cv=kf, scoring =

# print the score for each fold
print(np.sqrt(-scores).round(4))
print((sum(np.sqrt(-scores))/len(np.sqrt(-scores))).round(4))
print('RMSE score for Ridge model')

[0.3375 0.3352 0.3756 0.3229 0.3317]
0.3406
RMSE score for Ridge model
```

XGBoost

Bayesian hyperparameter optimisation for XGB

```
In [253...] from skopt.space import Real, Categorical, Integer
from skopt import BayesSearchCV
```

```
In [254...
from xgboost import XGBRegressor

model = XGBRegressor()

search_space = {
    'reg_lambda': Real(1e-10, 1e12, 'log-uniform'),
    'learning_rate': Real(0.005, 0.1),
    'n_estimators' : Integer(100, 5000),
    'max_depth' : Integer(2, 8),
    'subsample' : Real(0.5, 1.0),
    'colsample_bytree' : Real(0.25, 1.0),
}

xgb_opt = BayesSearchCV(model, search_space, cv = 5, n_iter= 8, scoring =
```

```
In [255...
#Thinking about making this as a warning.

print('Warning: Optimising and fitting XGBoost will take at least 40 minutes')
optimise = str(input('Would you like to optimise and fit XGBoost? Answer \

if optimise.lower() == 'c':
    model = xgb_opt.fit(X_train, y_train)
    xgb_opt.best_params_
```

Warning: Optimising and fitting XGBoost will take at least 40 minutes.

```
In [256...
xgb_opt.best_params_
```

```
Out[256... OrderedDict([('colsample_bytree', 0.6585666723543374),
    ('learning_rate', 0.02100437221055008),
    ('max_depth', 4),
    ('n_estimators', 2583),
    ('reg_lambda', 1.2994805179302653e-08),
    ('subsample', 0.6757498243587916)])
```

```
In [257...
#Why run the same 40 minute optimiser over and over? These are the best pa

xgboost = xgb_opt.best_estimator_
```

```
In [258...
score = cross_val_score(xgboost, X_val, y_val, cv=kf, scoring = 'neg_mean_s
print("xgboost: {:.4f} ({:.4f})".format(score.mean(), score.std()))
print(np.sqrt(-score)) #scores for each fold
```


```
xgboost: -0.0979 (0.0122)
[0.28633611 0.3105886 0.34594802 0.31012361 0.30812311]
```

```
In [259...
print(np.mean(np.sqrt(-score)))
```

```
0.3122238920801711
```


In [260...

```
from statlearning import plot_feature_importance
plot_feature_importance(xgb_opt.best_estimator_, labels=features)
plt.show()
```



In [280...

```
testxgb = xgboost.fit(fullxtrain, fullytrain)
```

In [282...

```
#predictions = yjp.inverse_transform(enet_cv.predict(test).reshape(-1,1))
predictions = np.round(np.exp(testxgb.predict(test)),2)
```

In [283...

```
my_submission = pd.DataFrame({'id': test.index, 'price': predictions.ravel})
my_submission.to_csv('submission.csv', index=False)
```

LightGBM

In [263...

```
import lightgbm as lgb
```

In [264...

```

model = lgb.LGBMRegressor(objective='regression')

search_space = {
    'learning_rate': Real(0.005,0.1), # uniform distribution between 0.005
    'n_estimators' : Integer(100, 2501), # discrete uniform distribution b
    'num_leaves' : Integer(2, 65), # discrete uniform distribution between
    'subsample' : Real(0.5, 1), # uniform distribution between 0.5 and 1
}

lbst = BayesSearchCV(model, search_space, n_iter = 32, cv = 5, n_jobs=-1,
lbst.fit(X_train, y_train)

print('Best parameters found by Bayes search:', lbst.best_params_, '\n')

```

```

Best parameters found by Bayes search: OrderedDict([('learning_rate', 0.088
38008753245956), ('n_estimators', 963), ('num_leaves', 4), ('subsample', 0.
7854148780763127)])

```

In [265...

```
lgb = lbst.best_estimator_
```

In [266...

```

score = cross_val_score(lgb, X_val, y_val, cv=kf, scoring = 'neg_mean_squa
print("LightGBM: {:.4f} ({:.4f})".format(score.mean(), score.std()))
print(np.sqrt(-score)) #scores for each fold

```

```

LightGBM: -0.1008 (0.0141)
[0.29353705 0.3197337 0.35638825 0.30437453 0.31011024]

```

In [267...

```
print(np.mean(np.sqrt(-score)))
```


```
0.31682875376343217
```

In [268...

```

from statlearning import plot_feature_importance
plot_feature_importance(lgb, labels=features)
plt.show()

```



```
In [269... #predictions = np.round(yj.inverse_transform(enet_cv.predict(test)).reshape(predictions = np.round(np.exp(lgb.predict(test)),2)
```

```
In [270... my_submission = pd.DataFrame({'id': test.index, 'price': predictions})
my_submission.to_csv('submission.csv', index=False)
```

Model Stacking

```
In [271... stack_train = X_train.values
stack_val = X_val.values
stack_y_val = y_val.values
stack_y_train = y_train.values
```

```
In [272... from mlxtend.regressor import StackingCVRegressor

stack = StackingCVRegressor(regressors=[lgb, enet, xgboost], meta_regressor
                             cv=5, use_features_in_secondary=True)
#modelstack = stack.fit(stack_train, stack_y_train)
```

```
In [274... score = cross_val_score(stack, stack_val, stack_y_val, cv=kf, scoring = 'neg
print("stack: {:.4f} ({:.4f})".format(score.mean(), score.std()))
print(np.sqrt(-score)) #scores for each fold
```

```
stack: -0.0962 (0.0141)
[0.28539588 0.31069724 0.35016403 0.2986285  0.30231729]
```

```
In [275... print(np.round(np.sqrt(-score),4))
```

[0.2854 0.3107 0.3502 0.2986 0.3023]

```
In [276... np.mean(np.sqrt(-score))
```

```
Out[276... 0.30944058692793924
```

```
In [277... teststack = stack.fit(fullxtrain, fullytrain)
```


```
In [278... stacktest=test.values
predictions = np.round(np.exp(teststack.predict(stacktest)),2)
```

```
In [279... my_submission = pd.DataFrame({'id': test.index, 'price': predictions.ravel()})
my_submission.to_csv('submission.csv', index=False)
```

Data Mining

```
In [366... train['logrevenue_per_month'] = np.log(train['revenue_per_month'])
```

```
In [367... from statlearning import plot_dists
plot_dists(train[['revenue_per_month', 'logrevenue_per_month']])
plt.show()
```



```
In [332... #These are the values for the data mining
```

```
In [368... features = train.columns.values
```

```
In [369... features = np.delete(features,np.where(features=='priceYJ'))
features = np.delete(features,np.where(features=='logprice'))
```

```
In [370... features = np.delete(features,np.where(features=='revenue_per_month'))
features = np.delete(features,np.where(features=='logrevenue_per_month'))
```

```
In [371... from sklearn.model_selection import train_test_split
X_tr_mine, X_val_mine, y_tr_mine, y_val_mine = train_test_split(train[featu
```

```
In [372... from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer

imp = IterativeImputer()
imp.fit(X_tr_mine[imputequantvars])
X_tr_mine[imputequantvars] = imp.transform(X_tr_mine[imputequantvars])
X_val_mine[imputequantvars] = imp.transform(X_val_mine[imputequantvars])
```

Should we be fitting the imputer to the test?

```
In [373... #Fixing variable skewness
```

```
In [374... skews = ['square_feet','weekly_discount','minimum_nights','minimum_minimum
```

```
In [375... yj = PowerTransformer(method='yeo-johnson')

X_tr_mine[skews] = yj.fit_transform(X_tr_mine[skews])
X_val_mine[skews] = yj.transform(X_val_mine[skews])
```

Outliers

```
In [376... quantitativevariables = []
for i in features:
    if train.dtypes[i] == int or train.dtypes[i] == float and i!='revenue_
        quantitativevariables.append(i)
```

```
In [377... X_tr_mine = X_tr_mine.fillna(0)
X_val_mine = X_val_mine.fillna(0)
```

```

In [378... from sklearn.linear_model import LinearRegression

ols = LinearRegression()
ols.fit(X_tr_mine[quantitativevariables], y_tr_mine)

y_fitted = ols.predict(X_tr_mine[quantitativevariables])
e = y_train - y_fitted

#e.reset_index(drop=True,inplace=True)

outliers = e.abs().sort_values().iloc[-10:] # retrieves the indexes for the

In [379... outlierlist = outliers.index.values

In [380... for i in outlierlist:
    X_tr_mine['outlier'] = np.where(X_tr_mine.index==i,1,0)

In [381... features = abs(train.corr()['logrevenue_per_month']).sort_values(ascending=

In [382... features=np.delete(features,np.where(features=='priceYJ'))
features=np.delete(features,np.where(features=='revenue_per_month'))
features=np.delete(features,np.where(features=='logprice')) #price is still

In [383... badfeatures = ['Nespresso machine', 'Natural gas barbeque', 'Sonos sound sy
#These features don't exist in the test and so will be deleted.
for i in badfeatures:
    features=np.delete(features,np.where(features==i)) # is causing it to
morebadfeatures = ['zhima_selfie', 'sent_id', 'reviews', 'manual_online
for i in morebadfeatures:
    features=np.delete(features,np.where(features==i)) # is causing it to

In [384... from sklearn.preprocessing import RobustScaler
scaler = RobustScaler() #scaling

X_train_scaled = scaler.fit_transform(X_tr_mine[features])
X_tr_mine = pd.DataFrame(X_train_scaled)

X_val_scaled = scaler.transform(X_val_mine[features])
X_val_mine = pd.DataFrame(X_val_scaled)

In [385... from sklearn.linear_model import ElasticNetCV
enet_cv = ElasticNetCV(l1_ratio=[0.01,0.05,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8
enet_cv.fit(X_tr_mine,np.ravel(y_tr_mine))

```

```
Out[385... ElasticNetCV(cv=5,
                        l1_ratio=[0.01, 0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8,
                                0.9,
                                0.95, 0.99, 1])
```

```
In [386... enet_cv.l1_ratio_
```

```
Out[386... 1.0
```

```
In [387... enet_cv.alpha_
```

```
Out[387... 0.004718791204316187
```

```
In [388... from sklearn.linear_model import ElasticNet
enetmine = ElasticNet(alpha = enet_cv.alpha_, l1_ratio = enet_cv.l1_ratio_
```

```
In [389... from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
kf=KFold(5)
scores = cross_val_score(enetmine, X_val_mine, y_val_mine, cv=kf, scoring =

# print the score for each fold
print(np.sqrt(-scores).round(4))
print((sum(np.sqrt(-scores))/len(np.sqrt(-scores))).round(4))
print('RMSE score for Elastic Net model')

[0.8116 0.7935 0.6941 0.7883 0.8011]
0.7777
RMSE score for Elastic Net model
```

```
In [390... enetmine.fit(X_tr_mine,y_tr_mine)
```

```
Out[390... ElasticNet(alpha=0.004718791204316187, l1_ratio=1.0)
```

```
In [395... coeffs = enetmine.coef_
```

```
In [396... d = {'Coefficient': coeffs.round(4)}

coeffmatrix = pd.DataFrame(data=d,index=features)
```

```
In [397... coeffmatrix.sort_values(by='Coefficient',ascending=False)
```

Out [397...

	Coefficient
host_acceptance_rate	0.9741
reviews_per_month	0.4042
host_response_rate	0.3783
price	0.2116
accommodates	0.1893
...	...
cleaning_fee_perc_blank	-0.1234
minimum_minimum_nights	-0.1458
Family/kid friendly	-0.1494
Internet	-0.1539
room_type_Private room	-0.3606

317 rows × 1 columns

In [398...

```
coeffmatrix.to_csv('coeffmatrix.csv', index=True)
```